

Natural Language Processing: Short Glossary

Token	The smallest unit of text processed by a language model.
Embedding	A numerical vector that represents the meaning of a token.
Attention	A mechanism that weighs which words are relevant to one another.
Transformer	A neural-network architecture built around attention layers.
Prompt	The instruction or text supplied to a language model.
Inference	The process of using a trained model to produce an output.
Hallucination	A fluent response that is unsupported or factually incorrect.
Fine-tuning	Additional training that adapts a model to a narrower task.